

A Multi-Framework Static Audit Tool for LLM System Prompts

Ali Murtaza Bhutto

Independent Researcher

ORCID: 0009-0007-2787-943X

Email: alibhutto101112@gmail.com

Abstract—We describe `ai-governance-checker`, a small Python tool that audits an LLM system prompt against three governance frameworks at once: the OWASP Top 10 for LLM Applications 2025 revision [1], the NIST AI Risk Management Framework 1.0 [2], and the EU AI Act high-risk-system checklist [3]. Each rule cites the specific framework section it implements, so a finding is traceable back to the authoritative text rather than to a private rulebook. An optional LLM-as-a-judge step [4] forwards the prompt to either Anthropic Claude or a local Ollama model for a structured second opinion, with three-step graceful degradation when no provider is reachable. On a labelled 30-prompt corpus (10 safe / 10 risky / 10 mixed) the rules-only path achieves micro-averaged accuracy = 0.9667, precision = 0.9406, recall = 0.9856, $F_1 = 0.9626$ across 16 framework section IDs (480 labelled cells). Eighty-nine pytest cases cover every rule’s positive and negative branches, the judge module’s three fallback paths, the aggregator, and the reporter. The repository is released under MIT and the eval is fully reproducible offline.

Index Terms—AI governance, OWASP LLM Top 10, NIST AI RMF, EU AI Act, LLM-as-a-judge, system prompt audit, compliance, prompt injection

I. INTRODUCTION

The deployment of LLM-based products at scale has outpaced the governance instruments meant to constrain them. Three frameworks now articulate the bulk of the practitioner’s compliance surface: the OWASP Top 10 for LLM Applications, refreshed in 2025 to add *System Prompt Leakage* and *Vector and Embedding Weaknesses* as new top-level categories [1]; the NIST AI Risk Management Framework 1.0 with its four cross-cutting functions GOVERN/MAP/MEASURE/MANAGE [2]; and the EU AI Act, Regulation (EU) 2024/1689, which makes specific obligations on transparency, human oversight, robustness, and record-keeping legally binding for high-risk AI systems in the European market [3].

A team preparing a release rarely has time to manually walk a system prompt through all three. Worse, the obvious threat surface in such prompts — prompt-injection vectors [5] including the indirect variant of Greshake et al. [6] — has been studied for several years but is still routinely missed in code review.

This paper describes `ai-governance-checker`, a small CLI that runs three rule packs simultaneously and produces a Markdown plus JSON audit report grouped by framework section. The contribution is not novel cryptography

or novel LLM safety theory; the contribution is a reproducible operational artefact that treats the three frameworks as parallel views on the same prompt and reports back with the framework-section citation visible on every finding.

Contributions. (1) Three rule packs implementing the parts of OWASP LLM Top 10 (2025), NIST AI RMF 1.0, and the EU AI Act that are inferable from a static system prompt — 20 rules in total, each tied to its published framework id. (2) An LLM-as-judge module with three-step graceful fallback (Anthropic → Ollama → DISABLED) so the tool stays useful when no API key is present. (3) A 30-prompt labelled eval corpus and an offline, deterministic eval harness that reports real measured precision / recall / F_1 per framework section. (4) A 90-test pytest suite and a hardened DevSecOps build (Bandit, pip-audit, Semgrep) so the tool can be vendored into a defender’s environment without first paying down a supply-chain debt.

II. RELATED WORK

The closest prior work is the LLM-as-a-judge methodology of Zheng et al. [4], which our optional judge module adopts directly. Where they evaluate model responses against each other, we evaluate the system prompt against a fixed governance checklist; the judging mechanism is the same, the substrate is not. We deliberately measure only the rule-pack path because a deterministic eval is reproducible offline; the judge path is supported but not benchmarked in this release, mirroring the honest stance our companion credential-exposure project takes for its keyed-API path.

For prompt-injection specifically, Perez and Ribeiro [5] were the first to catalogue the “Ignore Previous Prompt” family of attacks; the LLM01:2025 rule in our pack looks for the dual — explicit *role-lock* language that defends against those attacks. Greshake et al. [6] extended the threat to indirect injection via retrieved documents and tool outputs, which is what the LLM08:2025 *citation discipline* rule operationalises.

For governance instruments, the official texts of the OWASP top-10 2025 PDF, NIST AI 100-1, and Regulation (EU) 2024/1689 are the primary references for our rule packs. Each rule’s `framework_id` attribute is exactly the section identifier used in those documents, so a reviewer can cross-check a finding against the authoritative wording without first decoding our private rule names.

III. SYSTEM ARCHITECTURE




Fig. 1. Pipeline. Three rule packs and an optional judge feed an aggregator that emits paired Markdown and JSON reports.

A. Three rule packs, one report

The three frameworks are parallel views, not a hierarchy. The OWASP LLM Top 10 is written for security engineers; the NIST AI RMF is written for governance boards; the EU AI Act is written for regulators. A finding citing LLM01:2025 belongs in a security-engineering review; one citing GOVERN-1.1 belongs on a governance-board agenda; one citing EU-AIA-Art-13 ends up in a CE-mark conformity file. Folding all three into a single rulebook would lose this audit trail. We instead keep the three packs separate and have the reporter group findings by framework in the Markdown output.

B. LLM-as-judge with graceful fallback

The judge module follows a three-step decision: if ANTHROPIC_API_KEY is present, call the Claude Messages API; else if a local Ollama daemon is reachable, call its chat endpoint; else return JudgeStatus.DISABLED and let the report run in rules-only mode. Every branch is exercised by a respx-mocked pytest case (test_judge_uses_anthropic_when_key_present, test_judge_uses_ollama_when_anthropic_absent, test_judge_disabled_when_no_provider) so the absence of a key remains a tested code path rather than an untested conditional.

C. Severity bands

We use the five-level qualitative scale info / low / medium / high / critical, loosely aligned with the qualitative band in NIST SP 800-30 Rev. 1. The severity assignments are by category, not by case-by-case judgement: LLM02:2025 secret detections fire CRITICAL because a literal credential in a system prompt is unambiguously broken; LLM10:2025 unbounded-consumption findings fire LOW because they are operational rather than security-critical.

IV. EVALUATION

A. Test set

The labelled 30-prompt corpus in eval/synthetic_prompts.json contains 10 deliberately well-defended prompts (the *safe* class),

10 prompts that exercise specific framework violations (the *risky* class), and 10 partially-defended real-world-shaped prompts (the *mixed* class). Each prompt carries an expected_violations list naming the framework section IDs that should fire. We exclude the four always-fire INFO IDs (LLM03:2025, LLM04:2025 baseline, MANAGE-2.3, EU-AIA-Art-10) from the metric universe, leaving 16 discriminating IDs.

B. Measured numbers

TABLE I
RULES-ONLY EVAL, 30 PROMPTS $\times$ 16 FRAMEWORK IDS = 480 LABELLED CELLS, 2026-05-06.

Metric	Value
n prompts	30
n framework IDs	16
True positives	206
True negatives	258
False positives	13
False negatives	3
Accuracy	0.9667
Precision	0.9406
Recall	0.9856
F_1	0.9626
Per-ID F_1 mean	0.9614
Per-ID F_1 min	0.8000
Per-ID F_1 max	1.0000
Eval wall-clock	< 1 s

C. Test suite

The repository ships 90 pytest tests covering every rule’s positive and negative branches, the judge module’s three fallback paths under respx-mocked HTTP, the aggregator’s severity roll-up, and the reporter’s Markdown/JSON pair. Total wall-clock under one second.

D. Honest residuals

Thirteen false positives and three false negatives remain. They are not bugs we chose to suppress; they are real signal:

- *Partially-defended “safe” prompts.* Several have a clear role-lock and refusal clause but lack an explicit output-sanitisation directive (LLM05:2025) or an explainability clause (MEASURE-2.8). The rule fires correctly; the prompt could be tightened.
- *Phrasing close to but not exactly matching* one of the five robustness defence patterns. Tightening the regex further would over-fit to our labels rather than improve real coverage. We deliberately stopped tuning at this residual.

E. Honest negatives

We deliberately do *not* benchmark the LLM-as-judge module. Adding a paid-API axis would compromise the offline-only, deterministic property of the rules-only eval. The judge path is fully tested under respx mocks but its numbers against real provider models are out of scope for this release.

V. SECURITY POSTURE

The repository’s CI runs three independent scanners on every commit: Bandit (medium-and-above severity gate), pip-audit (transitive vulnerability scan against the OSV and PyPA advisory feeds), and Semgrep (p/python + p/security-audit rule packs). All three are green at the release tag with zero findings and zero suppressions.

VI. LIMITATIONS

Static-only. The tool sees the system prompt and nothing else. It cannot evidence training-data integrity (LLM04, EU AI Act Article 10), supply-chain controls (LLM03), or operational incident response (MANAGE-2.3). For these the rule pack emits an INFO finding pointing the reviewer at the parts of the deployment the rule pack cannot see.

Pattern-based. Rules are regex matchers, deliberately conservative. A prompt that uses an unusual idiom for a defence (“the assistant must defer to a senior reviewer when uncertain”) may be flagged for an oversight rule even though the defence is present in spirit. The reviewer is the final adjudicator; the rule pack is an aid.

Three frameworks, not all. ISO/IEC 23894:2023, the UK AISI inspect harness, and Singapore’s AI Verify are not yet encoded. Adding them is a matter of writing further rule packs in the same `CheckResult` interface; the architecture deliberately makes this additive.

VII. CONCLUSION

LLM governance tooling does not need new theory. What it needs is honest framework citations, graceful degradation when API keys are absent, and a small reproducible eval that demonstrates the rule pack actually does what its README claims. The 90-test, 20-rule, three-framework, MIT-licensed pipeline described here is one such instance.

REFERENCES

- [1] OWASP Foundation, “OWASP Top 10 for Large Language Model Applications, version 2025,” *OWASP GenAI Security Project*, Nov. 2024. [Online]. Available: <https://genai.owasp.org/llm-top-10/>
- [2] National Institute of Standards and Technology, “Artificial Intelligence Risk Management Framework (AI RMF 1.0),” NIST AI 100-1, Jan. 2023. [Online]. Available: <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- [3] European Parliament and Council, “Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act),” *Official Journal of the European Union*, L series, 12 July 2024.
- [4] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica, “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena,” in *Advances in Neural Information Processing Systems 36 (NeurIPS ’23)*, New Orleans, LA, USA, Dec. 2023.
- [5] F. Perez and I. Ribeiro, “Ignore Previous Prompt: Attack Techniques for Language Models,” in *NeurIPS 2022 ML Safety Workshop*, New Orleans, LA, USA, Dec. 2022.
- [6] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not What You’ve Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection,” in *Proc. 16th ACM Workshop on Artificial Intelligence and Security (AISec ’23)*, Copenhagen, Denmark, Nov. 2023, pp. 79–90.

Reproducibility colophon. The repository is at <https://github.com/thunderstornX/ai-governance-checker>. The eval re-runs in well under a second (python eval/run_eval.py). Pinned dependencies are listed in `requirements.txt`; the pytest suite (90 tests) runs in under a second on commodity hardware.